

curriculum vitae of

Pankaj Garkoti

✉ garkotipankaj@gmail.com
☎ +91 87663 35252

EDUCATION

2019 – 2022

B.Sc. Physics (H.)

DELHI UNIVERSITY, NEW DELHI

TECHNICAL SKILLS

- **Backend** Python, Go, TypeScript, FastAPI, Gin, Express, Django, async services, long running tasks
- **Frontend** Vue.js, React, Next.js, Tailwind, Vanilla JS + HTML + CSS, Chrome extensions
- **AI** LLM pipelines, fine tuning, retrieval, agents, OCR, memory, real-time voice
- **Cloud** AWS, GCP, Terraform, Docker, CI/CD
- **Integrations** Google Workspace, Slack, Jira, Notion, Asana, WhatsApp, anything that can be integrated to anything else TBH
- **Technical Leadership** system design & architecture, product research, UX, hiring and mentoring, managing engineering teams, being an early stage CTO

CAREER SNAPSHOT

- **Senior Software Engineer** — Jouser (DeepLaw), 2025 - 2026
- **Co-founder & CTO** — Mavex AI, 2023 - 2025
- **Co-founder & CTO** — Hear Us, 2022 - 2023

EXPERIENCE & BUILD LOG

Oct. 2025 – Jan 2026

Jouser Limited, British Virgin Islands

SENIOR SOFTWARE ENGINEER, DEEPLAW

First engineer at DeepLaw, one of three parallel AI product teams at Jouser. Built the product end-to-end from scratch in under 3 months before the company shut down when its sole investor pulled out.

- Architected the full contract analysis pipeline: parallelized OCR + extraction agents for risk factors and obligations, a party-resolution engine for legal entity disambiguation, and auto-scheduling of renewals to Google Calendar.
- Set up fully automated, auto-scaling infra on GCP with Terraform. Brought end-to-end analysis runtime under 50 seconds for documents up to 100 pages.

Apr. 2023 – May 2025

Mavex AI, Gurugram

CO-FOUNDER, CTO

An AI productivity platform for knowledge workers, raised \$100k+ across two rounds.

- Led all technical vision and execution across multiple product iterations — from a task-to-calendar blocking system, to a multi-agent hub with dynamically rendered UIs, to Mavy, an AI scheduling assistant CC'd directly into email threads.
- Designed what was likely one of the first delegation-based multi-agent architectures in production, before agentic tooling became mainstream.
- Managed a team of engineers, interns, and designers; owned the full stack including integrations with Google Calendar, Gmail, Notion, Asana, and Jira.

Jan. 2022 – Mar. 2023	HearUs, New Delhi	CO-FOUNDER, CTO
	A B2C mental health platform that pivoted twice and re-incorporated as Mavex AI after raising an angel round.	
	• Built a Flutter app for anonymous peer support, a Slack-based employee mood tracking tool (Joy), and a GPT-3-powered WhatsApp bot for scheduling — my first LLM-powered application. • Each pivot carried forward prior learnings; the technical and operational foundation became Mavex AI.	
Oct. 2021 – Apr. 2022	BCAS, Delhi University	RESEARCH APPRENTICE
	• Undertook the calculation of electron-molecule collision cross-sections for various elemental gases under the guidance of Dr. Anand Bharadvaja. • Presented the results in a research paper published in MDPI Journal — <i>Atoms</i>.	

PROJECTS

- **trie**
trie is a symbol-level codebase indexing layer that exposes an MCP server interface (**trie_grep**, **trie_read**, **trie_trace**), enabling AI coding agents to navigate a codebase by intent rather than brute-force file search. In benchmarks on the same question, agents with trie used 3 times less tool calls. The index bootstraps itself in the project and is incrementally kept in sync with the source code. In active development.
- **daily-notes.nvim**
daily-notes.nvim is a lightweight Neovim plugin that replicates Obsidian's markdown template based daily note functionality inside Neovim. The plugin by default carries over the contents of previous day's note, headers that should not be copied can be specified separately. The plugin is highly configurable and allows custom datetime formats and paths etc.
- **ChainFactory**
ChainFactory is a declarative system for building multi-step LLM workflows in a YAML-like syntax. Chains are composed of links whose outputs feed into subsequent inputs, with sequential and parallel execution, auto-generated prompts from purpose statements, and type-safe structured output via Pydantic models. Ideal for split-map-reduce style inference pipelines.
- **Computer Reinvention**
My personal open-source lab exploring how LLMs can reshape the way developers read, navigate, and reason about code. All work is public and built in the open.

PUBLICATIONS

- Garkoti, P.; Luthra, M.; Goswami, K.; Bharadvaja, A.; Baluja, K.L. The Binary-Encounter-Bethe Model for Computation of Singly Differential Cross Sections Due to Electron-Impact Ionization. *Atoms* 2022, 10, 60. *atoms*

REFERENCES

- Xuan Ma (CEO, Joustier Limited) (mxuan105@gmail.com)
- Saurabh Pandey (CEO, Mavex AI) (saurabh@mavex.ai)
- Dr. Anand Bhardvaja (Professor, University of Delhi) (anand.bharadvaja@bcas.du.ac.in)